

Thread Safety



DEVELOPMENTOR
DEVELOPING PEOPLE WHO DEVELOP SOFTWARE

Agenda

- **Highlight issues with multi threaded programming**
- **Introduce thread synchronization primitives**
- **Introduce thread safe collections**



Need for Synchronization

- **Creating threads is easy**
- **When threads share data problems can occur**
 - Inconsistent reads
 - State corruption
- **Synchronization fixes these problems, but potentially creates a new problem**
 - Over synchronization reduces scalability
- **Lots of techniques to implement synchronization**
 - Each have cost and benefit
- **Developers role is to write an application that scales and is thread safe, by selecting the best synchronization technique**



Simple Increment

- **Two threads**
 - Sharing an instance of **Counter**.
 - Both are calling **Increment** 1000 times
- **Question**
 - What is the value of count after both threads have completed?

```
public class Counter
{
    protected int count;

    public virtual void Increment()
    {
        count++;
    }
    public int Value { get { return count; } }
}
```



Simple Increment, NOT Atomic

- Even a simple `count++` is not an atomic operation.
 - Multiple CPU instructions that could be interweaved.
- Consider the possible execution below of two threads (T0, T1)
 - Assuming `count=0` at the start
 - At the end of execution i would be 1 and not the desired 2.
- If two threads don't attempt to increment count at the same time not a problem. Spotting these kind of errors is hard

time
↓

```
T0: MOV R0, count
```

```
T0: ADD R0, 1
```

```
T0: MOV count, R0
```

```
T1: MOV R0, count
```

```
T1: ADD R0, 1
```

```
T1: MOV count, R0
```



Interlocked

- **Modern CPU's expose special instruction set to perform various operations atomically**
 - Cost more than non atomic variants.
- **Access to these instructions via `Interlocked` class**
 - `Interlocked.Increment`
 - `Interlocked.Decrement`
 - `Interlocked.Add`
 - `Interlocked.Exchange`, `Interlocked.CompareAndExchange`
 - Useful for building Spin locks

```
public class InterlockedCounter : Counter
{
    public override void Increment()
    {
        // Atomic count++
        Interlocked.Increment(ref count);
    }
}
```



Multi step state transition

- **What happens if**
 - Thread A is **inside ReceivePayment**
 - Thread B is **inside NetWorth**
- **Can Interlocked help ?**

```
class SmallBusiness {  
    decimal Cash = 0;  
    decimal Receivables = 1000;  
  
    public void ReceivePayment(decimal amount) {  
        Cash += amount;  
        Receivables -= amount;  
    }  
  
    public decimal NetWorth {  
        get { return Cash + Receivables; }  
    }  
}
```



Sequential access

- **To fix the problem**
 - Sequentialise access to the object state
- **How**
 - Each instance of a reference type has a Monitor
 - CLR guarantees that only one thread can own the monitor
 - If a thread can't acquire the monitor it enters a wait state
 - When the monitor is available it is woken up and proceeds
- **Critical areas of code can therefore be protected by using a monitor.**



Monitor based solution

- Only one thread in any **critical region** at any point in time

```
private object _lock = new object();

public void ReceivePayment(decimal amount)
{
    Monitor.Enter(_lock);
    Cash += amount;
    Receivables -= amount;
    Monitor.Exit(_lock);
}
```

Could be an issue with exceptions

```
public decimal NetWorth
{
    get {
        Monitor.Enter(_lock);
        try { return Cash + Receivables; }
        finally { Monitor.Exit(_lock); }
    }
}
```

Deals better with exceptions



Lock keyword

- **Enter, try, finally , Exit common pattern**
 - C# language offers **lock** keyword to assist
 - Compiler emits try, finally logic
- **Use of Monitor.Enter and **lock** can lead to deadlocks**
 - Prefer Monitor.TryEnter which takes a timeout
- **Avoid using lock(this) and lock(typeof(X))**
 - Less control over objects use for synchronization.
 - Prefer creation of object for sole purpose of synchronization

```
public void ReceivePayment(decimal amount) {  
    lock (_lock)  
    {  
        Cash += amount;  
        Receivables -= amount;  
    }  
}
```



High Read to Write Ratio

- **Monitor provides mutual exclusion behaviour**
 - Excluding readers and writers
- **Thread safety not an issue if all threads read.**
- **Better throughput may be achieved with a Synchronization primitive that ensures**
 - There can be Many Readers, Zero Writer
 - Or One Writer, Zero Readers
- **This is known as a ReaderWriterLock**
 - .NET 3.5 and above prefer ReaderWriterLockSlim
 - Pre 3.5, ReaderWriterLock
 - Not well implemented, can result in writer being denied access for long periods of time.
- **Often used for caching, where most of the time is spent reading with occasional updates.**



Reader Writer Lock

```
public class SimpleCache
{
    private Dictionary<int, string> cache=
        new Dictionary<int, string>();
    private ReaderWriterLockSlim _lock = new ReaderWriterLockSlim();

    public string Get(int key) {
        _lock.EnterReadLock();
        try { return cache[key]; }
        finally { _lock.ExitReadLock(); }
    }

    public void Set(int key, string val)
    {
        _lock.EnterWriteLock();
        try { cache.Add(key, val) }
        finally { _lock.ExitWriteLock(); }
    }
}
```

Many threads can
can read from the
cache

When one thread
has the write lock
no other thread
can obtain read or
write lock.



Concurrent Collections

- **Standard Collections not thread safe**
 - List<T>, Dictionary<K, V>, Queue<T>, Stack<T>
- **.NET 4 introduces Concurrent variants**
 - ConcurrentQueue<T>
 - ConcurrentStack<T>
 - ConcurrentDictionary<K, V>
 - ConcurrentBag<T>
- **Designed to**
 - When possible to run lock free.
 - Operations don't block, TryXXXX



Example of Concurrent Collection

- **Dictionary.Add** throws exception if key already present
 - To prevent exception check if key is present, if not add item
 - Checking for key and then adding is two steps would require locking in a multithreaded environment
- **ConcurrentDictionary.TryAdd** will return false if key already exists. Negating the need for a lock

```
public class SimpleCache{  
    private ConcurrentDictionary<int, string> cache = new  
        ConcurrentDictionary<int, string>();  
  
    public string Get(int key) {  
        return cache[key];  
    }  
    public void Set(int key, string val) {  
        cache.TryAdd(key, val);  
    }  
}
```




Rendezvous


- **Co-operating tasks sometime need to synchronize with each other before proceeding.**
 - Exchange results.
 - Wait for all tasks to initialise before commencing.
- **Barrier**

```
Barrier phase = new Barrier(3);
```


```
phase.SignalAndWait();
```



```
phase.SignalAndWait();
```



```
phase.SignalAndWait();
```

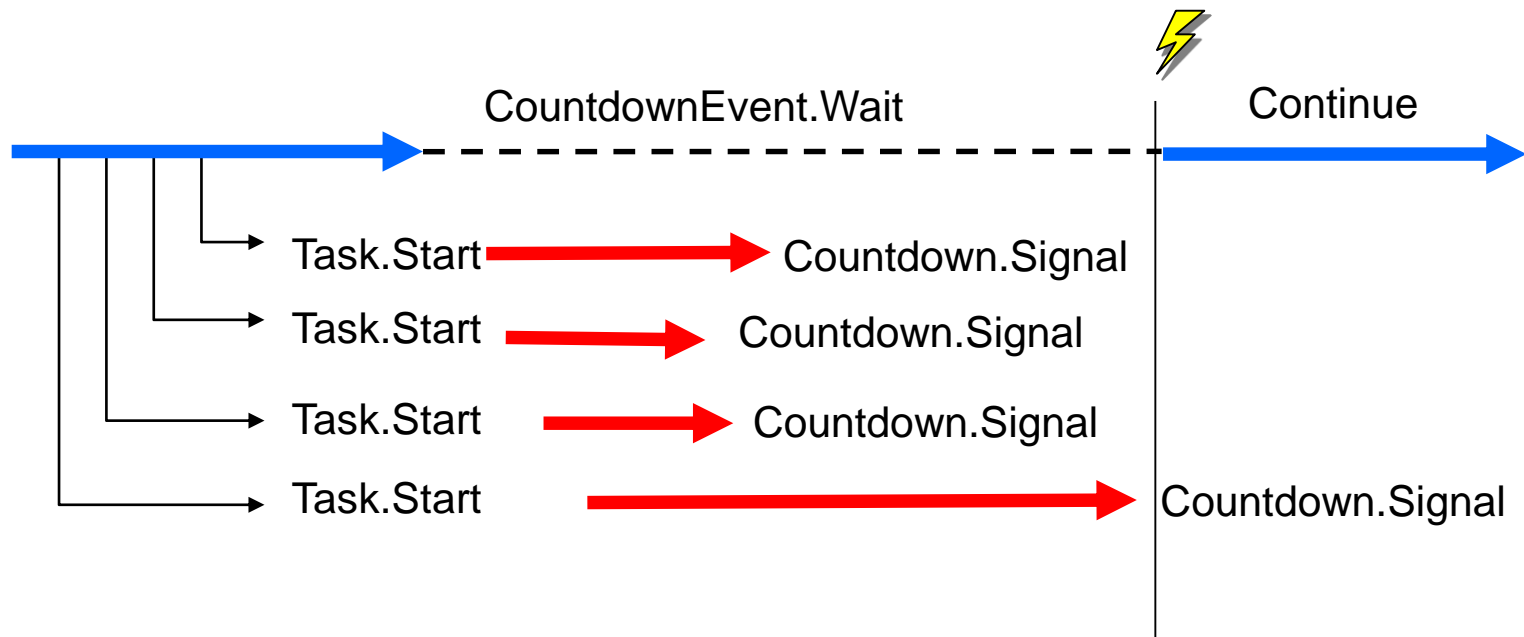


Threads block until all three have called SignalAndWait



Countdown Event

- **Primary** task initiates N other **tasks** and wishes to wait for them all to have completed a series of steps.
- Creates an instance of a **CountdownEvent** initialised to the number of sub tasks.



Synchronization across app domains

- **Managed synchronization primitives only allow synchronization inside a single app domain**
- **How to control access to a shared file ?**
 - Requires Kernel based synchronization
- **Kernel synchronization can be achieved via managed wrappers**
 - Mutex
 - Semaphore
 - AutoResetEvent
 - ManualResetEvent
- **These synchronization primitives are orders of magnitude more expensive than managed ones**



Summary

- **A variety of ways to perform synchronization, the skill is picking the correct one**
- **Concurrent collections make it simpler to write efficient thread safe code**
- **Only use kernel synchronization primitives when absolutely necessary**
- **Analyse code and imagine worse possible race conditions**

