

Parallel Programming



DEVELOPMENTOR


DEVELOPING PEOPLE WHO DEVELOP SOFTWARE

Objectives

- **Why program in parallel ?**
- **Introduce Parallel Extensions**
 - Parallel.Do
 - Parallel.For
 - Parallel.ForEach
- **Introduce PLINQ Parallel Linq**



“Free Lunch is over” Herb Sutter, Dr Dobb’s Journal March 2005



Why the need for parallel programming

- **Doesn't MS Word run fast enough ?**
- **For a given time frame X**
 - Data volume increases
 - More detail in games
 - Intellisense
 - Better answers
 - Speech recognition
 - Route planning
 - Word gets better at grammar checking



Setting Expectations

- **Amdahl's Law**

- If only $n\%$ of the compute can be parallelised at best you can get an $n\%$ reduction in time taken.
- Example
 - 10 hours of compute of which 1 hour cannot be parallelised.
 - Minimum time to run is 1 hour + Time taken for parallel part, thus you can only ever achieve a maximum of 10x speed up.

- **Not even the most perfect parallel algorithms scale linearly**

- Scheduling overhead
- Combining results overhead



Types of Parallelism, Coarse Grained

- **Coarse grain parallelism**
 - Large tasks, little scheduling or communication overhead
 - Each service request executed in parallel, parallelism realised by many independent requests.
 - DoX,DoY,DoZ and combine results. X,Y,Z large independent operations
 - Applications
 - Web Page rendering
 - Servicing multiple web requests
 - Calculating the payroll



Types of Parallelism, Fine Grained

- **Fine grain parallelism**

- A single activity broken down into a series of small co-operating parallel tasks.
 - $f(n=0...X)$ can possibly be performed by many tasks in parallel for given values of n
 - Input \Rightarrow Step 1 \Rightarrow Step 2 \Rightarrow Step 3 \Rightarrow Output
- Applications
 - Sensor processing
 - Route planning
 - Graphical layout
 - Fractal generation (Trade Show Demo)



Moving to a Parallel Mind Set

- **Identify areas of code that could benefit from concurrency**
 - Keep in mind Amdahl's law
- **Adapt algorithms and data structures to assist concurrency**
 - Minimise shared state
 - Lock free data structures
- **Utilise class `System.Threading.Parallel`**
 - Built on top of the Task library



Fork/Join Pattern

- **Task based parallelism**
 - Book Hotel, Book Car, Book Flight
 - Sum X + Sum Y + Sum Z

```
Parallel.Invoke( BookHotel , BookCar , BookFlight );
```

```
// Won't get here until all three tasks are complete
```

```
Task<int> sumX = Task.Factory.StartNew(SumX);
```

```
Task<int> sumY = Task.Factory.StartNew(SumY);
```

```
Task<int> sumZ = Task.Factory.StartNew(SumZ);
```

```
int total = sumX.Result+ sumY.Result+ sumZ.Result;
```



Parallel Loops

- **Loops can be successfully parallelised IF**
 - The order of execution is not important.
 - Each iteration is independent of each other.
 - Each Iteration has sufficient work to compensate for scheduling overhead.



Parallel.For

- **Replaces a conventional for loop with a parallel loop**
 - Restrictions
 - Only integer range.
 - Increments only by 1.
- **Loop body supplied by a delegate.**
- **Will the output of both these pieces code be the same ?**

```
for (int i = 0; i < 10; i++)  
{  
    Console.WriteLine(i);  
}
```

```
Parallel.For(0, 10, i =>  
{  
    Console.WriteLine(i);  
});
```



Parallel.For Mechanics

- **Parallel.For utilises multiple tasks**
 - The number of created tasks will vary based on number of cores and monitored throughput.
 - A single task may execute many loop iterations.
 - Each iteration of the loop results in a delegate invocation
- **Since loop iterations run in parallel**
 - All the normal thread safety issues come into play
 - Heavy use of synchronization can negate the full benefit.



Parallel Aggregation

- Would the following `Parallel.For` be a suitable replacement ?

```
double pi = 1;  
double multiplier = -1;
```

```
for (int i = 3; i < 100000000; i+=2) {  
    pi += multiplier * (1.0 / (double)i);  
    multiplier = multiplier * -1;  
}
```

```
pi *= 4.0;
```

```
Parallel.For(0, 100000000 / 2, i =>  
{  
    pi += multiplier * (1.0 / ((double)3 + i * 2.0));  
    multiplier *= -1;  
});
```



Thread Safe Parallel.For

- Utilising Monitor based synchronization around updating shared state would create an unnecessary bottleneck.
- Ideally we want each task to have no shared state and can run freely.
- **Parallel.For** allows each Task to have a piece of local storage.
 - Utilise local storage to build a partial result.
 - Supply a delegate to be run when task is first used to initialise local result.
 - Supply a delegate to be run when the task is no longer used to combine its local result with a shared global result
 - **WARNING**...still need to synchronize for this part.



Thread Safe Parallel.For

Initialise local state

local state

```
double pi = 1;
object combineLock = new object();

Parallel.For(0, 100000000 / 2, () => 0.0, (i, loopState, localPi) =>
{
    double multiplier = i % 2 == 0 ? -1 : 1 ;
    localPi += multiplier * (1.0 / ((double)3 + i * 2.0));

    return localPi;
}, localPi => {
    lock (combineLock)
        pi += localPi;
    } ));

pi *= 4.0;
```

Return update local state for next iteration of this task

Run this piece of code when task is no longer required to combine local result with global value. NOTE Synchronization still required.



Parallel.ForEach

- **Similar to Parallel.For but parallelises the consumption of a IEnumerable<T>**
- **IEnumerator<T> is not thread safe so access is guarded**
 - Initially one item is taken per spawned task
 - Attempts to learn the optimal amount to take per task, thus reducing contention.

```
List<int> values = new List<int> { 1, 2, 3, 4, 5 };  
Parallel.ForEach(values, i =>  
{  
    Console.WriteLine(i);  
});
```



Parallel Loop Options

- **Parallel Options**
 - CancellationToken
 - Supply your own Cancellation Token, to allow client termination of the loop.
 - Max Degree of Parallelism
 - The maximum number of tasks will run concurrently for the loop.
 - Task Scheduler
 - Select an alternative scheduler.



Alternative Parallel.For

- **Parallel.For is very limited**
 - Only allow int's
 - No steps, just ++
- **Consider using Parallel.ForEach with an iterator method**

```
Parallel.ForEach(Range(1.0, 2.0, 0.1), i =>
{
    Console.WriteLine(i);
});

private static IEnumerable<double> Range(
    double start, double end, double step) {
    for ( double i = start; i < end; i+= step){
        yield return i;
    }
}
```



My Parallel.For doesn't scale

- **Small loop bodies often don't parallelise**
 - Either they don't scale as expected
 - Or in some cases are slower than sequential version
- **Overhead of task management and delegate invocation cost**
 - As number of cores increases this will become less noticeable.
- **Parallel loops are typically best achieved when you have an inner sequential loop.**

```
for( ;; )  
{  
    for( ;; )  
    {  
  
    }  
}
```

Parallelise outer
loop to ensure
sufficient work
per delegate
invocation

```
Parallel.For( )  
{  
    for( ;; )  
    {  
  
    }  
}
```



Creating Outer loops

- A single for loop, can be replaced by an outer and inner loop.

```
Parallel.For (0,5000,i=>
{
    Do(i);
} );
```

```
Parallel.For(0, 50, outerIndex =>
{
    for (int innerIndex = 0; innerIndex < 100; innerIndex++){
        int i = innerIndex + 100 * outerIndex;
        Do(i);
    }
});
```



Parallel.ForEach outer loop

- **Refactoring single loops to outer inner loops can be tedious**
 - Consider using a iterator method to generate a stream of ranges to represent the bounds of the inner for loops

```
Parallel.ForEach(new Range(0, 5000).CreateSubRanges(50), r =>
{
    for (int i = r.Start; i < r.End; i++)
    {
        Do(i)
    }
});
```



Outer loop generation

```
public struct Range    {
    public int Start;
    public int End;

    public IEnumerable<Range> CreateSubRanges(int nRanges)    {

        int subRangeStart = Start;
        int subRangeStep = ((End - Start)+1) / nRanges;
        while (nRanges > 1) {
            yield return new Range()
                { Start = subRangeStart,
                  End = subRangeStart + subRangeStep-1 };
            subRangeStart += subRangeStep;
            nRanges--;
        }

        yield return new Range() { Start = subRangeStart, End = End };
    }
}
```



PLinq

- The **AsParallel ()** extension method on **IEnumerable<T>**, results in using parallel versions of **Where,Select** etc.
 - Works with Linq to Objects and Linq to XML

Utilise Parallel Extension methods
Instead of Enumerable

```
var values = new List<int> { ... };

var results = from val in values.AsParallel()
              where val % 2 == 0
              select val;

foreach (int result in results)
{
    Console.WriteLine(result);
}
```



PLinq Partitioning

- **PLinq queries start with partitioning**
 - How much data and which bits of data to give each task.
- **Range Partitioning**
 - When the source has an index able interface
 - Queries the `IEnumerable<T>` to see if its an `T[]` or `IList<T>`
- **Chunk Partitioning**
 - Tasks request a chunk of items, the size of chunks typically grow over time.
 - Can produce better load balancing.
 - Used for non index able sources.
- **Utilising Range Partitioning can increase performance.**



ForAll

- **Parallel version of List<T>.ForEach**
 - ForAll takes an action delegate
 - Action is executed in parallel.

```
var values = new List<int> { ... };  
  
values.AsParallel().ForAll(i =>  
{  
    if (i % 2 == 0)  
    {  
        Console.WriteLine(i);  
    }  
});
```



Summary

- **Is the Free Lunch back ?**
 - Utilising Parallel.XXX may bring back a moderate free lunch
 - Still need to understand what is happening under the covers if you want to scale, and have a gut busting lunch.
- **Benchmark code sequentially before attempting to parallelise**
 - Benchmark on a range of cores.
- **Maximise the use of local storage to reduce contention and simplify thread safety.**

