

Inside the Garbage Collector



DEVELOPMENTMENTOR
DEVELOPING PEOPLE WHO DEVELOP SOFTWARE

Agenda

- **Applications and Resources**
- **The Managed Heap**
- **Mark and Compact**
- **Generations**
- **Non-Memory Resources**
- **Finalization**
- **Hindering the GC**
- **Helping the GC**

Applications and Resources



- **Applications use many different kinds of resources when executing**
 - Memory
 - File handles
 - Database connections
 - Sockets
- **.NET provides infrastructure and patterns for managing an applications resources**

Memory and the Managed Heap

- **.NET has specific functionality to automate memory management**
 - Applications request memory using new keyword
 - Memory is allocated in area called the **Managed Heap**
- **Runtime cleans up memory when necessary using Garbage Collection (GC)**
 - Only managed heap is garbage collected. Unmanaged and stack allocated memory unaffected
 - Garbage Collector removes objects no longer in use

How the Garbage Collector Works

- Applications run as if memory is infinite
- GC is invoked when resource threshold breached
- Assumes all objects are garbage
- Collects those **not in use**



What Objects are in Use?

- **The garbage collector does not collect objects in use**
 - How does it know what is in use and what is not?
- **GC has concept of objects being reachable**
 - Objects that could be accessed from code that is yet to run
- **GC walks objects graph under live roots**
 - Non-null static references
 - Local variables on the stack that are still in use
 - A few more exotic kinds, e.g. the native code interop infrastructure
- **Any object reachable from a live root survives the GC**

Spot the Live Roots

- What references are live roots when the GC runs?

```
static object o1 = new object();

void RunTest()
{
    object o2 = new object();
    object o3 = new object();

    GC.Collect();

    o1 = null;

    GC.Collect();

    ProcessObject(o3);

    GC.Collect();
}
```


What is garbage here?

What is garbage here?

What is garbage here?

GC Optimization - Generations

- **Examining the entire heap every GC too expensive**
 - Optimization: if object survived last GC then probably still alive
 - By default only checks the most recently allocated objects
- **.NET GC is Generational**
 - Three generations
 - All objects allocated into gen0
 - Any object surviving a collection is promoted to next generation



The Implications of Generations

- **Relative cost of GC**
 - Gen0 is very cheap to collect
 - Gen1 is more expensive to collect than gen0
 - Gen2 can be very expensive to collect
- **Healthy GC**
 - Order of magnitude between number of collections of each successive generation regarded as healthy
 - $\text{gen0} : \text{gen1} : \text{gen2} \Rightarrow 100 : 10 : 1$
- **Generally don't try to second guess the GC**

GC Optimization – GC Styles

- **Three styles of GC**
 - Non-concurrent
 - Concurrent
 - Asynchronous

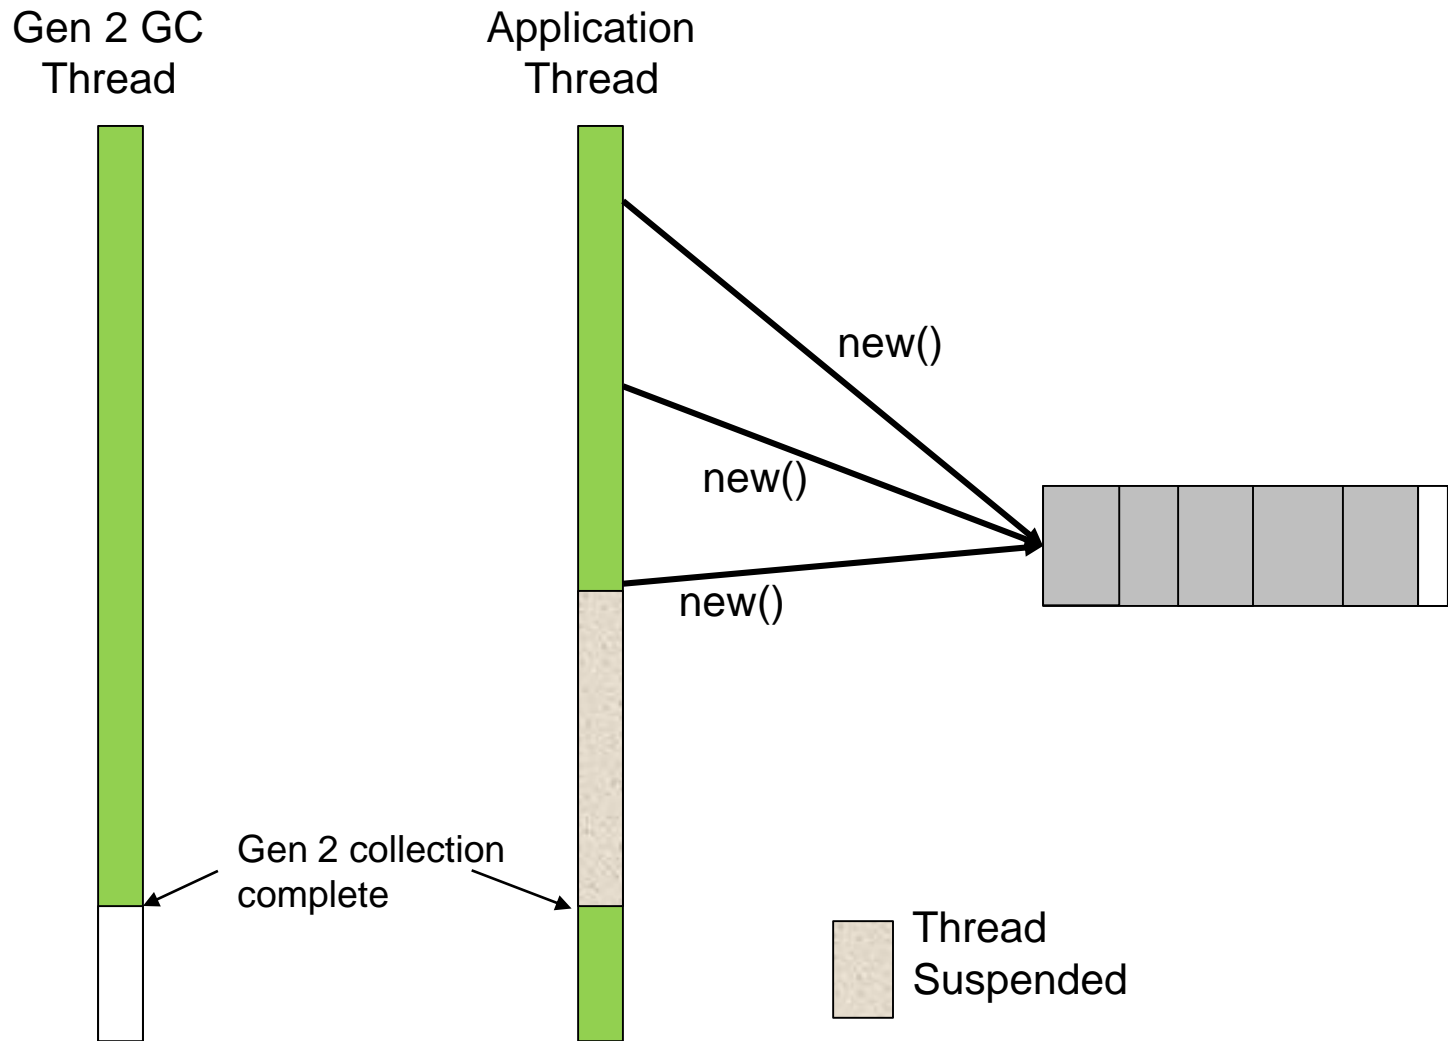
Concurrent and Non-Concurrent

- **Non-concurrent**
 - GC starts and runs through to completion without interruption
 - Always used for Gen 0 and Gen 1 collects
- **Concurrent**
 - Application threads interleaved with collection
 - Keeps application responsive
 - Only used for some Gen 2 collects
 - Can continue application execution until another GC required

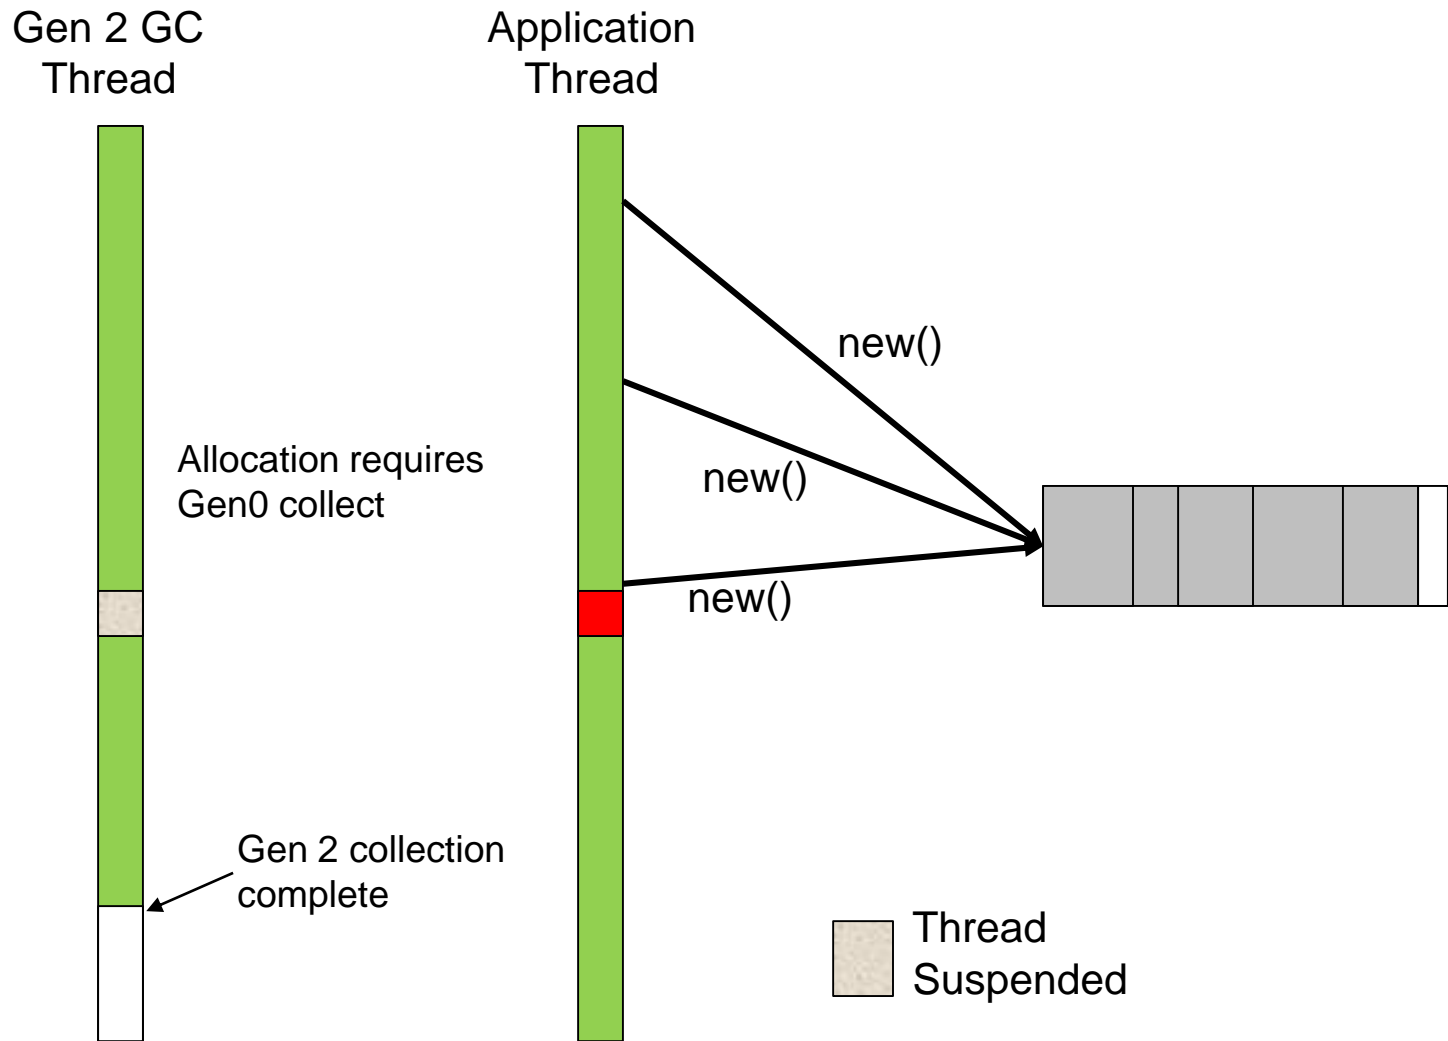
Async

- **Asynchronous**
 - Introduced in version 4.0
 - Similar to concurrent but can perform gen 0 and gen 1 collects during a gen 2 collect

Concurrent GC



Background GC



GC Modes

- **There are three GC modes**
 - Workstation Concurrent (default)
 - Workstation Non-Concurrent
 - Server
- **Controlled by config flags**

```
<configuration>
  <runtime>
    <gcServer enabled="true"/>
    <gcConcurrent enabled="false"/>
  </runtime>
</configuration>
```

Workstation Concurrent

- **Designed to keep UI applications interactive**
 - UI thread paused as short a time as possible
- **Gen 0 and Gen 1 use non-concurrent GC**
 - UI thread pause time would be insignificant
- **Gen 2 can use concurrent**
 - Gen 2 can be expensive
 - Keeps UI thread pumping the message queue
- **Gen 2 can use asynchronous**
 - Means UI thread stays free-running even if another GC is required
 - Only available on .NET 4.0 and later

Workstation non-concurrent

- **Whole GC is non-concurrent**
 - Optimized for non-UI code
 - GC kidnaps allocating thread
 - Fewer CPU cycles than workstation concurrent

Server GC

- **Tuned for multiple processors**
 - Each processor gets own block of managed heap for allocation / GC called arena
 - Allows lock free allocation across multiple cores
- **Always non-concurrent**
 - GC takes place on all processors at the same time

GC Optimization – Self Tuning

- **GC behavior dependent on your application**
 - Uses heuristics to work out efficient strategy
 - If GC of a generation is effective then performs more
 - If GC of a generation reclaims little memory then performs less
- **GC self tuning normally produces optimal results**
 - Best strategy for your application
- **Avoid GC.Collect**
 - Disrupts GC statistics and so self tuning less effective

Large Objects

- **Moving large objects too expensive**
 - Large object over 85000 bytes
- **Allocated in Large Object Heap**

Memory is not the only Resource

- **Only memory management is automated**
 - Runtime does not know when use of other resources is complete
- **Types that control non-memory resources follow standard pattern**
 - Implement IDisposable
 - Consumer calls Dispose when they are finished

```
public interface IDisposable
{
    void Dispose();
}
```

When Should You Implement IDisposable?

- **Most classes do not need to implement IDisposable**
 - Do not consume resources other than memory
- **Classes that directly control resources acquired via interop must implement IDisposable**
 - No other managed type can release those resources
- **Classes that aggregate other types that implement IDisposable should implement IDisposable themselves**
 - Provides a hook to Dispose aggregated objects

What if I don't Call Dispose?

- **Only an issue for types that directly control unmanaged resources**
 - FileStream
 - SqlConnection
 - Icon
- **Need last chance mechanism for freeing up resources**
 - No one else can free up resources
- **Runtime has concept of Finalization**
 - Very unlikely that your classes will need a finalizer

How Finalization Works

- **Types that support finalization override Finalize virtual method on System.Object**
 - C# syntax looks like a C++ destructor

```
class SharedMemory
{
    ~SharedMemory()
    {
    }
}
```

- **Runtime sees type has finalizer when instantiated**
 - Puts object in finalization queue

Finalization and GC

- **When GC decides to collect object it sees it's finalizable**
 - Moves the object on to freachable queue
 - freachable queue acts as live root (finalization **reachable**)
- **Separate finalization thread processes freachable queue**
 - Calls finalizer and removes object from queue
 - Object can be collected next time GC looks at it

Consequences of Finalization

- **Finalizable objects will always survive one GC**
 - Always promoted to gen1 or gen2
 - Consequently more expensive to collect
- **Timing of finalizer execution non-deterministic**
 - After GC notices it is garbage
 - When Finalization thread gets round to it
- **Limited functionality in finalizer**
 - Should not call other objects as may already have been finalized
- **Finalizable classes should always implement IDisposable**
 - Can [suppress finalization](#) in Dispose

```
public void Dispose()  
{  
    GC.SuppressFinalize(this);  
}
```

Hindering and Helping the GC

- **Three main issues for GC**
 - Unnecessary promotion
 - Reference heavy data structures
 - Pinning

Unnecessary Promotion

- **Want to minimise Gen2 collections**
 - Ideally all Gen2 objects live for ever
 - Not realistic
- **Promote then die is worst scenario**
 - Badly tuned caches
 - Finalizers
 - “ad-hoc” caching
 - Not releasing dead objects

Unnecessary Promotion

- **Tune cache timeouts**
 - Drop cache in Gen1
 - Make it long lived to make the hit worthwhile
- **Remove unnecessary finalizers**
 - If absolutely required refactor code into SafeHandle

```
public abstract class SafeHandle : CriticalFinalizerObject,  
                                   IDisposable  
{  
    ...  
}
```

“Ad-hoc” Caching

“Probably need this object again soon so I’ll hang on to it”

- **Consider WeakReference**
 - Object accessible unless GC’d

```
WeakReference wr = new WeakReference(new object());  
  
object temp = wr.Target;  
if(temp == null)  
{  
    Console.WriteLine("Collected");  
}  
else  
{  
    Console.WriteLine("Still alive");  
}
```

Hanging on to Dead Objects

- **Set member and static references to null when done**
 - GC can clean up as soon as possible
- **Especially important when replacing object**
 - Be careful of threading issues

```
class ReferenceData
{
    XElement data;


    public void ReloadData(string filename)
    {
        data = null;
        data = XElement.Load(filename);
    }
}
```

Reference Heavy Structures

- **Mark phase made harder by many references**
 - GC has to chase down embedded object graphs
- **Think of revised or optimised ways of representing data**
 - Arrays
 - Adjacency Matrices
- **Example: Thread Pool Queue before 4.0**


Thread Pool Queue Before 4.0

- **Linked list of work items**
 - Root is never garbage
 - Lots of work = lots of references



Thread Pool Queue After 4.0

- **Linked list of arrays of work items**
 - Lots of work = few references



Pinning

- **Pinning means GC can't move object**
 - Required in interop scenarios
- **Consider inflating large buffers to be allocated in LOH**
 - Pinning irrelevant in LOH
 - Reuse LOH buffers

Summary

- **GC automates memory reclamation**
- **GC is heavily optimized**
- **Release mode GC much more aggressive than debug**
- **Non memory resources freed up using IDisposable pattern**
- **Finalization exists as a last resort for framework classes. You should not need it in your code**
- **You can help the GC**
- **You can hinder the GC**